

NAVIIN RAJ A/L THEIVARAJAN

Kuala Lumpur, MY | +6016-355 2145 | naviin2373@gmail.com | [linkedin.com/in/xxufnaviin](https://www.linkedin.com/in/xxufnaviin) | github.com/xxufnaviin

ABOUT ME

[Qualifications](#)

Computer Science student specializing in Data Analytics, currently focusing on data engineering and infrastructure. Interested in distributed computing, resource optimization, and building efficient compute environments. Gaining hands-on experience with cluster management, orchestration tools, and cloud-based infrastructure.

TECHNICAL SKILLS

Programming Languages: Python, C/C++, Java, SQL, R, GoLang, YAML, JavaScript, HTML, CSS

Data Engineering: Apache Airflow, Apache Spark, ETL (Extract, Transform, Load), Distributed Systems, Redis, PostgreSQL

DevOps/Cloud: CI/CD, Docker, AWS (ECR, EC2, ECS, S3), Google (Cloud Storage, Compute Engine), Terraform, Bash

AI/ML: Scikit-Learn, Matplotlib, Seaborn, Jupyter, Agentic AI, MCP Server, LLM

Tools: VSCode, MSSQL, Github, Gitlab, AWS Management Console, Google Cloud Platform, Datadog, Postman

Soft Skills: Leadership, Problem-Solver, Empathetic, Motivated, Communication Skills, Time Management

EXPERIENCE

Data Engineer Intern – Grab

Jul 2025 - Nov 2025

- Assisted in migrating Spark jobs into a unified platform by packaging job dependencies and functions as semantic-versioned Python packages (Poetry, pyproject.toml) and automating deployment with GitLab CI/CD
- Refactored legacy Airflow integrations by replacing custom APIs with official Apache Airflow APIs in Go, supported by Postman testing and debugging
- Enabled automatic cluster engine selection using data/artifacts stored in AWS S3, enhanced Spark-on-Kubernetes with improved annotations and job submission logic, and implemented regex-based token masking on backend gateway services to secure sensitive data across the unified platform
- Migrated and centralized Docker images to AWS ECR with container scanning to ensure vulnerability-free deployments, and replicated legacy workflows while automating downstream pipeline triggers via curl requests
- Executed large-scale code and dependency migrations using AI-assisted tooling, updating secret paths and dependencies to improve security and functionality

FINAL YEAR PROJECT

Logarda – Cloud Monitoring System | AWS

Apr 2026 – Present

Analytics Module | Python, Bash, Docker, CI/CD, ECR, EC2, PostgreSQL, Redis | [Github](#)

- Developed an **automated data pipeline** to ingest compute metrics and log messages from all users into **PostgreSQL**, while queuing error logs in **Redis** to support **scalable downstream processing**
- Containerized** the ingestion workflow with **Docker** and deployed it to **ECR**, running on **self-hosted EC2** as CRON jobs
- Currently developing **data processing pipeline** to prepare metrics data for **time-series forecasting models**
- Currently developing **LLM inference pipeline** supported by **RAG-based workflow** using **AWS** documentations
- Currently developing **API endpoints** for LLM inference on error log messages

Backend Module | Golang, PostgreSQL, Redis | [Github](#)

- Developed a scalable **backend** system supporting concurrent users, implementing **AES-GCM encryption** and **SHA-256 hashing** to enhance data security and integrity
- Currently developing **metrics stream worker** to update dashboard when user is connected
- Currently developing **error message inference worker** to monitor **Redis queue** for error messages
- Currently developing **API endpoints** to handle client-side requests

PROJECTS

Stocky AI - Decision Intelligence Agent | Python, CI/CD, AWS(ECR, ECS, S3), Docker | [Github](#) Apr 2026 - May 2026

- Designed the **overall system architecture**, integrating Agentic AI, MCP tools, and API services
- Implemented **Unit Test**, **Integration Test**, **E2E Test** using PyTest to validate code changes in **CI/CD** pipeline
- Deployed the **Dockerized** solution to **ECR**, maintained using **ECS** for high availability

AuroraFlow - Weather ETL | Spark, Airflow, Docker, GCS, GCE, Python | [Github](#) Aug 2025 - Nov 2025

- Developed a scalable **ETL pipeline** using **PySpark** and **Airflow** to process real-time weather data
- **Containerized** the workflow with **Docker** and deployed Airflow services on **Google Compute Engine**.
- Integrated **Google Cloud Storage (GCS)** for storage and downstream analytics

L-Atelier, AI Powered Trend Analysis Platform | Airflow, Google BigQuery, Python, FastAPI | [Github](#) Sep 2025

- Designed the **overall system architecture**, integrating data processing, ML components, and API services
- Built an **ETL pipeline** that ingested YouTube data, extracted MP3 audio, and transcribed it for downstream analysis
- Developed **FastAPI endpoints** for trend analytics retrieval and implemented **Airflow DAGs**.

ACHIEVEMENTS

1st Runner-Up in UM Hackathon 2026

- Developed Stocky AI - Decision Intelligence Agent, along with my teammates
- Contributed to the project by designing system architecture, writing test cases, and handling automated deployments

2nd Runner-Up in UM Datathon 2025

- Developed machine learning based cash flow forecasting with executive dashboards, along with my teammates
- Contributed to the project by creating executive dashboards using Power BI and pitching part of the solution

Finalist in UM Hackathon 2025

- Developed Grabby AI, an intelligent assistant designed to help Grab merchants, along with my teammates
- Contributed to the project by coordinating tasks and work delegations and handling Q&A during live pitching

LANGUAGES

Able to read, write and speak fluently in **English, Mandarin and Malay**.

EDUCATION

Asia Pacific University Bukit Jalil, KL

Bachelor of Science (Honours) in Computer Science, Specialism in Data Analytics Oct 2023 - Present

- CGPA 3.87
- Vice Chancellor's List in the Academic Year 2023/2024
- Vice Chancellor's List in the Academic Year 2024/2025