

NAVIIN RAJ A/L THEIVARAJAN

Kuala Lumpur, MY | +6016-355 2145 | naviin2373@gmail.com | [linkedin.com/in/xxufnaviin](https://www.linkedin.com/in/xxufnaviin) | github.com/xxufnaviin

ABOUT ME

Portfolio

Computer Science student specializing in Data Analytics, currently focusing on data engineering and infrastructure. Interested in distributed computing, resource optimization, and building efficient compute environments. Gaining hands-on experience with cluster management, orchestration tools, and cloud-based infrastructure.

TECHNICAL SKILLS

Programming Languages: Python, C/C++, Java, SQL, R, GoLang, YAML, JavaScript, HTML, CSS

Data Engineering: Apache Airflow, Apache Spark, Apache Iceberg, ETL, Redis, PostgreSQL, Hadoop, HDFS

DevOps/Cloud: CI/CD, Docker, AWS (ECR, EC2, ECS, S3), Google (Cloud Storage, Compute Engine), Terraform, Bash

AI/ML: Scikit-Learn, Matplotlib, Seaborn, Jupyter, Agentic AI, MCP Server, LLM, RAG, LangChain

Tools: VSCode, MSSQL, Github, Gitlab, AWS Management Console, Google Cloud Platform, Datadog, Postman

Soft Skills: Leadership, Problem-Solver, Empathetic, Motivated, Communication Skills, Time Management

EXPERIENCE

Data Engineer Intern – Grab

Jul 2025 - Nov 2025

- Assisted in migrating Spark jobs into a unified platform by **packaging job dependencies** and functions as semantic-versioned Python packages (Poetry, pyproject.toml) and automating deployment with **GitLab CI/CD**
- Refactored legacy Airflow integrations by **replacing custom APIs** with official Apache Airflow APIs in **Go**, supported by **Postman testing** and debugging
- Enabled **automatic cluster engine selection** using data/artifacts stored in **AWS S3**, enhanced Spark-on-Kubernetes with **improved annotations** and **job submission logic**, and implemented **regex-based token masking** on backend gateway services to secure sensitive data across the unified platform
- Migrated and **centralized Docker images** to **AWS ECR** with container scanning to ensure **vulnerability-free deployments**, and replicated legacy workflows while **automating downstream pipeline triggers** via curl requests
- Executed large-scale code and dependency migrations using **AI-assisted tooling**, updating secret paths and dependencies to improve security and functionality

FINAL YEAR PROJECT

Logarda – Cloud Monitoring System | AWS, PostgreSQL, Redis, EC2

Apr 2026 – Present

Analytics Module | Python, Bash, Docker, CI/CD, ECR | [Github](#)

- Developed an **automated data pipeline** to ingest compute metrics and log messages from all users into **PostgreSQL**, while queuing error logs in **Redis** to support **scalable downstream processing**
- Containerized** the ingestion workflow with **Docker** and deployed it to **ECR**, running on **self-hosted EC2** as CRON jobs
- Developed a **data processing pipeline** for preparing cloud metrics data for time-series forecasting.
- Built an **LLM inference pipeline** using a **Retrieval-Augmented Generation (RAG)** workflow with AWS documentation.
- Designed a **hybrid GRU + Linear Regression** forecasting model for resource metric prediction, achieving **5.56% sMAPE**.

Backend Module | Golang, WebSocket, REST API, Real-Time Systems | [Github](#)

- Developed a scalable **backend** system supporting concurrent users, implementing **AES-GCM encryption** and **SHA-256 hashing** to enhance data security and integrity
- Developed a **metrics streaming worker** to provide real-time dashboard updates through **event-driven data processing**.
- Built an error message inference worker to consume Redis queue messages and perform automated error analysis.
- Developed **REST API endpoints** to handle client requests and support cloud monitoring functionalities.
- Implemented a **WebSocket gateway** for real-time communication between connected clients and the backend server.

PROJECTS

Lakeura – Data Lakehouse AI Agent | Python, LangChain, Apache Iceberg, Hadoop | [Github](#) Sep 2026 - Oct 2026

- Developed an **AI data platform** that enables natural language interaction with raw **CSV/Parquet** datasets
- Built an **MCP-based** data layer using **PySpark** and **Apache Iceberg** for data onboarding, querying, and table management
- Integrated **LangChain, Groq, and MCP** to build an agentic workflow for automated data discovery and management

Stocky AI - Decision Intelligence Agent | Python, CI/CD, AWS(ECR, ECS, S3), Docker | [Github](#) Apr 2026 - May 2026

- Designed the **overall system architecture**, integrating Agentic AI, MCP tools, and API services
- Implemented **Unit Test, Integration Test, E2E Test** using PyTest to validate code changes in **CI/CD** pipeline
- Deployed the **Dockerized** solution to **ECR**, maintained using **ECS** for high availability

AuroraFlow - Weather ETL | Spark, Airflow, Docker, GCS, GCE, Python | [Github](#) Aug 2025 - Nov 2025

- Developed a scalable **ETL pipeline** using **PySpark** and **Airflow** to process real-time weather data
- **Containerized** the workflow with **Docker** and deployed Airflow services on **Google Compute Engine**.
- Integrated **Google Cloud Storage (GCS)** for storage and downstream analytics

ACHIEVEMENTS

1st Runner-Up in UM Hackathon 2026

- Developed **Stocky AI - Decision Intelligence Agent**, along with my teammates
- Contributed to the project by designing system architecture, writing test cases, and handling automated deployments

2nd Runner-Up in UM Datathon 2025

- Developed machine learning based cash flow forecasting with executive dashboards, along with my teammates
- Contributed to the project by creating executive dashboards using Power BI and pitching part of the solution

Finalist in UM Hackathon 2025

- Developed **Grabby AI**, an intelligent assistant designed to help Grab merchants, along with my teammates
- Contributed to the project by coordinating tasks and work delegations and handling Q&A during live pitching

LANGUAGES

Able to read, write and speak fluently in **English, Mandarin and Malay**.

EDUCATION

Asia Pacific University Bukit Jalil, KL

Bachelor of Science (Honours) in Computer Science, Specialism in Data Analytics Oct 2023 - Present

- CGPA 3.87
- Vice Chancellor's List in the Academic Year 2023/2024
- Vice Chancellor's List in the Academic Year 2024/2025

REFERENCE

Arun Ravi M V | Data Engineering Manager II

Email: arunravi.mv@grab.com

Company: Grab (MyTeksi Sdn. Bhd.)

Relationship: Internship Supervisor/Manager